

Assignment #1: Containers with Docker

The purpose of this assignment is to build and run a simple two-container stack using Docker and Docker Compose. One container will run a PostgreSQL database, while the other will host a lightweight Python application. The Python app will connect to the database, query a few rows, compute basic statistics, and then print and save the results. This exercise introduces you to the fundamentals of multi-container setups, including service networking, environment variables, and reproducible workflows. These are essential skills that you will continue to use and expand upon in the course project.

Learning goals:

- Run a database in a container and seed it with initial data.
- Connect from another container using service DNS, ports, and env vars.
- Query records and compute simple statistics.
- Package everything to run with one command.
- Use GitHub to manage your code and configuration.

What you will build:

- Service 1: PostgreSQL
 - Starts with a seeded table: `trips(id SERIAL PRIMARY KEY, city TEXT, minutes INT, fare NUMERIC(6,2))`
 - Data is inserted via an `init` script.
- Service 2: Python app
 - Connects to Postgres using `env` vars.
 - Runs a couple of read queries.
 - Computes simple stats.
 - Writes `out/summary.json` and prints a small table to `stdout`.

1. Installation Requirements

To complete this assignment, you must have the following tools installed and working on your machine:

- **Docker Desktop:** required to build images and run containers.
- **Docker Compose:** for defining and running multi-container applications.
- **GitHub Desktop (optional):** for version control and to manage your repository.
- **IDE:** such as VS Code, PyCharm, or any editor you prefer.

In addition:

- Make sure you can run basic Docker commands (docker build, docker run, docker ps) before you start.
- The database is only reachable on the internal Docker network, so no host port is used, and port 5432 does not need to be free. If you choose to add a mapping to the db service to connect from your host, make sure the port is free first.
- You should be comfortable creating and navigating directories, editing files, and running commands from the terminal.
- Create a **public** GitHub repository for this assignment, named ITCS6190-A1-<your-name>. DON'T add the instructor and the TAs as collaborators.

Suggested GitHub repo layout:

```
.
├─ app/           main.py, Dockerfile
├─ db/           init.sql, Dockerfile
├─ out/          summary.json (created at runtime)
├─ compose.yml
├─ Makefile
├─ .gitignore
└─ README.md
```

Add a **.gitignore** containing at least out/, __pycache__/, and .env, so that generated output and any local credentials stay out of the repository.

2. Setup the Database service container

Create the Dockerfile for the database: this file defines the database container. It starts from the official Postgres image and copies your initialization script into the special folder that Postgres checks on first startup. Any .sql files in that folder will be executed automatically, allowing you to preconfigure tables and seed data.

db/Dockerfile
FROM postgres:16
COPY init.sql /docker-entrypoint-initdb.d/

Create the initialization SQL script: this script sets up your schema and inserts a few rows of sample data. When the container runs for the first time, Postgres will automatically execute the script. Expand it later for your final tests.

db/init.sql (example)

```
CREATE TABLE trips (  
    id SERIAL PRIMARY KEY,  
    city TEXT NOT NULL,  
    minutes INT NOT NULL,  
    fare NUMERIC(6,2) NOT NULL  
);  
INSERT INTO trips (city, minutes, fare) VALUES  
(  
    'Charlotte', 12, 12.50),  
    ('Charlotte', 21, 20.00),  
    ('New York', 9, 10.90),  
    ('New York', 26, 27.10),  
    ('San Francisco', 11, 11.20),  
    ('San Francisco', 28, 29.30);
```

How it works: when the db container starts for the first time, Postgres automatically executes any scripts in `/docker-entrypoint-initdb.d/`. This ensures your database is initialized with the correct schema and seed data without manual setup.

3. Setup the Python app service container

Create the Dockerfile for the app: This file defines the Python service container. It starts from a lightweight Python image, installs `psycopg` (a PostgreSQL driver), sets the working directory, copies your application code into the container, and specifies the default command to run.

app/Dockerfile

```
FROM python:3.11-slim  
RUN pip install --no-cache-dir psycopg[binary]==3.1.19  
WORKDIR /app  
COPY main.py /app/  
CMD ["python", "main.py"]
```

Create the Python application: This script connects to the PostgreSQL database using environment variables passed in via `compose.yml`. It retries the connection until the database is ready, executes a few queries (total trips, average fare per city, top N longest trips), and outputs the results both to the console and as a JSON file in `/out/summary.json`. Complete the missing parts.

app/main.py

```
import os, sys, time, json  
import psycopg  
  
# Environment variables with defaults  
DB_HOST = os.getenv("DB_HOST", "db")  
DB_PORT = int(os.getenv("DB_PORT", "5432"))  
DB_USER = ...  
DB_PASS = ...  
DB_NAME = ...  
TOP_N = int(os.getenv("APP TOP N", "5"))
```

```
def connect_with_retry(retries=10, delay=2):
    last_err = None
    for _ in range(retries):
        try:
            conn = psycopg2.connect(
                host=DB_HOST,
                port=DB_PORT,
                user=DB_USER,
                password=DB_PASS,
                dbname=DB_NAME,
                connect_timeout=3,
            )
            return conn
        except Exception as e:
            last_err = e
            print("Waiting for database...", file=sys.stderr)
            time.sleep(delay)
    print("Failed to connect to Postgres:", last_err, file=sys.stderr)
    sys.exit(1)

def main():
    conn = connect_with_retry()
    with conn, conn.cursor() as cur:
        # Total number of trips
        cur.execute("...")
        total_trips = cur.fetchone()[0]

        # Average fare by city
        cur.execute("""
            ...
            ...
            ...
            """)
        by_city = [{"city": c, "avg_fare": float(a)} for (c, a) in cur.fetchall()]

        # Top N trips by DEFINE YOUR QUERY
        cur.execute("""
            ...
            """, (TOP_N,))
        top = ...

    summary = {
        "total_trips": int(total_trips),
        "avg_fare_by_city": by_city,
        "top_by_minutes": top
    }

    # Write to /out/summary.json
    os.makedirs("/out", exist_ok=True)
    with open("/out/summary.json", "w") as f:
        json.dump(summary, f, indent=2)

    # Print to stdout
    print("=== Summary ===")
    print(json.dumps(summary, indent=2))

if __name__ == "__main__":
    main()
```

4. Compose the two services together

Compose the two services: Now that you have created the database and Python app containers, the next step is to define how they work together using Docker Compose. The `compose.yml` file specifies both services, their environment variables, volumes, and dependencies. This configuration allows you to launch the entire stack with a single command, so the database and application run together seamlessly. Complete the missing parts of the file.

```
compose.yml
services:
  db:
    build:
      context: ./db
      dockerfile: Dockerfile
    environment:
      POSTGRES_USER: appuser
      POSTGRES_PASSWORD: secretpw
      POSTGRES_DB: appdb
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U appuser -d appdb"]
      interval: 3s
      timeout: 3s
      retries: 10
  app:
    build:
      context: ./app
      dockerfile: Dockerfile
    depends_on:
      db:
        condition: service_healthy
    environment:
      TODO: DB_HOST, DB_PORT, DB_USER, DB_PASS, DB_NAME
      APP_TOP_N: "10"
    volumes:
      - ./out:/out
```

Run the stack: Once the `compose.yml` file is set up, you can run both containers together with:

```
docker compose up --build
```

Note on credentials. The database user and password here are throwaway values used only on your machine, so committing them in `compose.yml` is acceptable for this assignment. In real projects, they belong in a `.env` file listed in `.gitignore` and referenced from Compose.

This command builds the images (if they are not already built), starts the PostgreSQL database service, and then runs the Python app. The app connects to the database, runs its queries, and writes the output to the `out/` folder. You should also see the JSON summary printed directly to your terminal.

5. One-command run

Write a Makefile: at this stage, you can simplify your workflow by using a Makefile for common Docker Compose commands. This ensures you can build, start, and clean up your containers with a single, easy-to-remember command. Also, your `README.md` should provide clear documentation so others (or you later) can run the stack without confusion.

```
Makefile
.DEFAULT_GOAL := all
COMPOSE=docker compose
build:
    $(COMPOSE) build app
up:
    $(COMPOSE) up --build
down:
    $(COMPOSE) down -v
clean: down
    rm -rf out && mkdir -p out
all: clean up
.PHONY: build up down clean all
```

Tip: Use tabs (not spaces) before each command.

Test run *make* in your project folder. Without `.DEFAULT_GOAL`, make would run only the first target (build); with it, make runs all, which cleans, builds the images, starts both services, and produces the output JSON. You can also run any target by name, for example make down.

6. Document the work

Document: every project must include a clear `README.md` file at the root of the repository. The README explains what your stack does and provides all the information someone else needs to build and run it without asking you questions. Think of it as your project’s “user guide”. It should include at least:

- What the stack does (2–3 sentences)
- Exact commands to run/stop
- Example output (copy/paste)
- Where outputs are written
- Troubleshooting (DB not ready, permission on `out/`, etc.)

7. Submission

Submit your final deliverables on Canvas. Each student must turn in:

- The URL of the GitHub repository containing your assignment code, Dockerfiles, `compose.yml`, and `README`.
- The repository must be **public** (don't add the instructor and TAs as collaborators), and must include a `.gitignore`.
- A 1-page PDF that includes:
 - The exact commands you used to build and run the stack.
 - A pasted summary from `stdout` (the JSON printed by the app).
 - The contents of `out/summary.json`.
 - A short reflection (3–5 sentences) describing what you learned and what you would improve.

You can export the PDF directly from your `README` or write it separately. Keep it concise, professional, and easy to read.

8. Grading rubric

- Correct multi-container setup (20 pts): Compose runs; services reachable; app reads from DB; summary written to `/out`.
- Data model and seed (15 pts): Proper table + init data via `init.sql`.
- App correctness (25 pts): Queries execute; basic stats computed; error handling for failures.
- Reproducibility (10 pts): One-command run; clear `README`.
- Code and organization (10 pts): clean repo structure following the layout above; small Docker image; `.gitignore` present; no real credentials committed.
- Reflection (20 pts): Clear, specific takeaways.

The 100 points above are scaled to the 2 points this assignment contributes to your final grade.